

# gcc-14: What's new and exciting for RISC-V

**Bay Area RISC-V Meetup, Santa Clara, CA**  
**May 23, 2024**

Vineet Gupta, Rivos

# Wait, why are we doing this?

Because why not :)

There seems to be a comprehensive resource out there already

<https://gcc.gnu.org/gcc-14/changes.html>

Well

*"If a picture is worth a thousand words, a patch is worth at least a thousand pictures". [Will Deacon on [lkml](#)]*

Share Interesting Technical tidbits in gcc-14

Share what Rivos Team has been up to in gcc software ecosystem

Drink some free Beer ;-)

# gcc-14 for RISC-V Highlights

## **Working Auto-Vectorization**

- Mature enough to run entire SPEC CPU2017 (SPECint as well as SPECfp)

## **Support for ever growing list of New ISA extensions**

- Zicond, Zfa to name a few

Improved code for Atomics

Pressure Sensitive Instruction Scheduling

Constant Synthesis Improvements

Initial Sign Extension Elimination Improvements

# gcc-14 RISC-V updates: Auto-Vectorization Works

## User writes

```
int abs(int i) {
    return i < 0 ? -i : i;
}

int sum_of_abs_diff (unsigned char *img1,
                    unsigned char *img2, int n)
{
    int sum = 0;
    for (int i = 0; i < n; i++)
        sum += abs (img1[i] - img2[i]);
    return sum;
}
```

## But Compiler needs to ...

```
for (int i = 0; i < N; i++) {
    // implicit type promotion
    int pix1 = (int) img1[i];
    int pix2 = (int) img2[i];
    // 32-bit sub (w/ sign-ext)
    int diff = pix1 - pix2;
    sum += abs (diff);
}
```

gcc 13.1  
-O2 -march=rv64gcv\_zbb  
-ftree-vectorize

```
abs:
    sraiw    a5,a0,31
    xor      a0,a0,a5
    subw     a0,a0,a5
    ret

sum_of_abs_diff:
    ble      a2,zero,.L6
    mv       a4,a0
    add      a2,a0,a2
    li       a0,0
.L5:
    lbu      a3,0(a1)
    lbu      a5,0(a4)
    addi     a4,a4,1
    addi     a1,a1,1
    subw     a5,a5,a3
    sraiw    a3,a5,31
    xor      a5,a3,a5
    subw     a5,a5,a3
    addw     a0,a5,a0
    bne      a2,a4,.L5
    ret

.L6:
    li       a0,0
    ret
```

gcc 14.1  
-O2 -march=rv64gcv\_zbb -ftree-vectorize

```
abs:
    negw     a5,a0
    max      a0,a5,a0
    ret

sum_of_abs_diff:
    ble      a2,zero,.L6
    vsetvli  a5,zero,e32,m1,ta,ma
    vmv.v.i  v2,0
.L5:
    vsetvli  a5,a2,e8,mf4,ta,ma
    vle8.v   v3,0(a1)
    vle8.v   v4,0(a0)
    sub      a2,a2,a5
    add      a0,a0,a5
    add      a1,a1,a5
    vwsbu.vv v1,v4,v3
    vsetvli  zero,zero,e16,mf2,tu,mu
    vmvlr.v  v3,v2
    vmslt.vi v0,v1,0
    vneg.v   v1,v1,v0.t
    vwadd.wv v2,v3,v1
    bne      a2,zero,.L5
    vsetvli  a5,zero,e32,m1,ta,ma
    li       a4,0
    vmv.s.x  v1,a4
    vredsum.vs v2,v2,v1
    vmv.x.s  a0,v2
    ret

.L6:
    li       a0,0
    ret
```

# gcc-14 RISC-V updates: Auto-Vectorization #2

Both SLP and Loop Vectorizers enabled in the middle end

Length Agnostic (VLA) and Length Specific (VLS) supported

- Length Agnostic reads VLEN at runtime
- Stripmine style loops

## Why all the fuss

- RVV ISA is special: Vector Insns predicated twice
  - VLEN mask and Regular v0 mask
- Needed new patterns in the middle-end:
  - WHILE\_LEN
- Initial patch for RVV support was a mere ~3M lines in ~1400 files
  - <https://gcc.gnu.org/pipermail/gcc-patches/2022-May/595903.html>

```
.../gcc.target/riscv/rvv/intrinsic/vwsub.c | 21607 ++++++
.../gcc.target/riscv/rvv/intrinsic/vwsubu.c | 21607 ++++++
.../gcc.target/riscv/rvv/intrinsic/vxor.c | 63631 ++++++
.../gcc.target/riscv/rvv/intrinsic/vzext.c | 10087 +++
.../gcc.target/riscv/rvv/stack/rvv-stack.exp | 47 +
.../rvv/stack/stack-check-alloca-scalar.c | 53 +
.../rvv/stack/stack-check-alloca-vector.c | 45 +
.../stack/stack-check-save-restore-scalar.c | 48 +
.../stack/stack-check-save-restore-vector.c | 62 +
.../riscv/rvv/stack/stack-check-scalar.c | 205 +
.../rvv/stack/stack-check-vararg-scalar.c | 33 +
.../riscv/rvv/stack/stack-check-vector_1.c | 277 +
.../riscv/rvv/stack/stack-check-vector_2.c | 141 +
1409 files changed, 3197907 insertions(+), 199 deletions(-)
create mode 100644 gcc/config/riscv/md-parser
create mode 100644 gcc/config/riscv/riscv-insert-vsetvl.cc
create mode 100644 gcc/config/riscv/riscv-vector-builtins-fur
```

# gcc-14 RISC-V updates: VSETVL Pass

## Backbone of Auto-Vec in the RISC-V backend

- Compiler starts off by adding a VSETVL insn for each Vector insn
- A dedicated RISC-V Pass then uses standard Lazy Code Motion (LCM) algorithm to
  - Eliminate redundant VSETVL insns
  - Fuse compatible VSETVL insns and ...
  - ... Hoist them outside the loop

```
#include "riscv_vector.h"

void foo1 (void * restrict in, void * restrict out, int n)
{
    for (int i = 0; i < n; i++) {
        vuint8mf8_t v = *(vuint8mf8_t*)(in + i);
        *(vuint8mf8_t*)(out + i) = v;
    }
}
```

```
-O2 -march=rv64gcv
--param=vsetvl-strategy=simple
```

```
foo1:
    ble     a2,zero,.L5
    add     a2,a0,a2
.L3:
    vsetvli a5,zero,e8,mf8,ta,ma
    vle8.v  v1,0(a0)
    vsetvli a5,zero,e8,mf8,ta,ma
    addi    a0,a0,1
    vse8.v  v1,0(a1)
    addi    a1,a1,1
    bne     a0,a2,.L3
.L5:
    ret
```

```
-O2 -march=rv64gcv
```

```
foo1:
    ble     a2,zero,.L1
    add     a2,a0,a2
    vsetvli a5,zero,e8,mf8,ta,ma
.L3:
    vle8.v  v24,0(a0)
    addi    a0,a0,1
    vse8.v  v24,0(a1)
    addi    a1,a1,1
    bne     a0,a2,.L3
.L1:
    ret
```

<https://godbolt.org/z/5a5GMc4fE>

# gcc-14 Updates: Constant Synthesis Improvements

## SPEC2017 Deepsjeng dynamic icounts increased in gcc 13.1 (regression)

- Top blocks involve Large constants
- RV ISA doesn't allow efficient encoding of Large constants
- But compiler are not supposed to make it any worse
- Compiler strategy changed over times
  - gcc 12 used constant pools
  - gcc-13 started synthesizing them but terrible codegen
  - gcc-14 relaxed register usage restriction and reuse same lo/hi 32-bit values

2023-05-09 c104ef4b5eb1 RISC-V: improve codegen for large constants with same 32-bit lo and hi parts [2]

2023-04-11 0530254413f8 riscv: relax splitter restrictions for creating pseudos

2023-03-01 7e52f4420ffb RISC-V: costs: miscomputed shiftadd\_cost triggering synth\_mult [PR/108987]

```
long long foo(void)
{
    return 0x01010101010101u11;
}
```

gcc 12.1: -O2 -march=rv64gc

```
f:
    lui    a5,%hi(.LC0)
    ld     a0,%lo(.LC0)(a5)
    ret

.LC0:
    .dword 72340172838076673
```

gcc 13.1: -O2 -march=rv64gc

```
f:
    li     a0,0x101_000
    addi   a0,a0,0x101
    slli   a0,a0,16
    addi   a0,a0,0x101
    slli   a0,a0,16
    addi   a0,a0,0x101
    ret
```

gcc 14.1: -O2 -march=rv64gc

```
f:
    li     a5,0x101_0000
    addi   a5,a5,0x101
    slli   a0,a5,32
    add    a0,a0,a5
    ret
```

<https://godbolt.org/z/hTTabeY4e>

# gcc-14 Updates: Sign Extension Elimination Improv #1

**Extraneous Sign Extension in RISC-V gcc codegen is a systematic problem.**

**ISA and the ABI are to blame, but again compilers can do better**

- Built into most basic of ALU insns: ADDW/ADDIW
  - signed 32-bit addition and final sign-extension into 64-bit reg
- RISC-V ABI mandates that 32-bit values (on rv64) are held as sign extended values in 64-bit wide reg.
  - Simplifies calling convention
  - Allows 64-bit comparison insn to work on 32-bit values

**Decent amount of work went in gcc-14, but lot more needs to done!**

2023-04-28 1966741378d5 RISC-V: Eliminate redundant zero extension of minu/maxu operands

2023-06-07 99bfd072e67 RISC-V: Eliminate extension after for \*w instructions

2023-10-24 fb4e2c1648ea RISC-V: elide unnecessary sign extend when expanding cmp\_and\_jump

2023-10-16 8eb9cdd14218 expr: don't clear SUBREG\_PROMOTED\_VAR\_P flag for a promoted subreg [target/111466]



# gcc-14 Updates: Sign Ext Elim Improv #2

## Interesting tidbit with the last change in last slide

2023-10-16 8eb9cdd14218 expr: don't clear SUBREG\_PROMOTED\_VAR\_P flag for a promoted subreg [target/111466]

### Just deletes a couple of lines

```
diff --git a/gcc/expr.cc b/gcc/expr.cc
index 308ddc09e631..d259c6e53385 100644
--- a/gcc/expr.cc
+++ b/gcc/expr.cc
@@ -9332,13 +9332,6 @@ expand_expr_real_2 (sepop ops, rtx target,
    op0 = expand_expr (treeop0, target, VOIDmode, modifier);
-   if (TYPE_UNSIGNED (TREE_TYPE (treeop0)) != unsignedp
-       && GET_CODE (op0) == SUBREG)
-     SUBREG_PROMOTED_VAR_P (op0) = 0;
-
    return REDUCE_BIT_FIELD (op0);
```

### Added back in 1994 !

(with empty changelog)

**1994-07-08 506980397227** (expand\_expr, case CONVERT\_EXPR): If changing signedness and we have a promoted SUBREG, clear the promotion flag.

On 10/5/23 08:56, Richard Kenner wrote:

Obviously, I have no recollection of that change at all. In July of 1994, I don't believe I was actively working on much in the way of ports, though I could be misremembering. My guess is that this change was to fix some bug, but I'm a bit mystified why I'd have batched so many different changes together in one ChangeLog entry like that. I think you're correct that it was most likely the Alpha that showed the bug.

# gcc-14 Updates: Sign Ext Elim: Glimpse of work to be done

## Redundant Extension Elimination REE Pass needs some love

- Can't handle "USE with more than One reaching DEFinition" or if they are in different BBs
- Can't leverage the ABI
  - Incoming "int" args are already sign-extended

```
/* PR/111467 */
int max(int a, int b)
{
    return (a > b) ? a : b;
}
```

gcc 14.1: -O2 -march=rv64gc

```
max(int, int):
    mv     a5,a0
    bge    a0,a1,.L4
    mv     a5,a1
.L4:
    sext.w a0,a5
    ret
```

clang: -O2 -march=rv64gc

```
max(int, int):
    blt     a1, a0, .LBB1_2
    mv     a0, a1
.LBB1_2:
    ret
```

<https://godbolt.org/z/qb3jvrnoe>

gcc -fdump-rtl-ree

```
# DEF 1
(insn 8 4 11 2 (set (reg:SI 15 a5 [orig:137 b ] [137])
  (reg:SI 11 a1 [orig:136 b ] [136]))) {movsi_internal}

(jump_insn 11 8 22 2 (set (pc)
  (if_then_else (ge (reg/v:DI 11 a1 [orig:136 b ] [136])
    (reg/v:DI 10 a0 [orig:135 a ] [135])))
  (label_ref 13)
  (pc))) "max.c":12:20 273 {*branchdi}
  (int_list:REG_BR_PROB 536870916 (nil))
-> 13)

# DEF 2
(insn 12 22 13 3 (set (reg:SI 15 a5 [orig:137 b ] [137])
  (reg:SI 10 a0 [orig:135 a ] [135]))) {movsi_internal}

(code_label 13 12 23 4 2 (nil) [1 uses])

# USE: Multiple reaching DEFs
(insn 19 14 20 4 (set (reg/i:DI 10 a0)
  sign_extend:DI (reg:SI 15 a5 [orig:137 b ] [137])))
  {extendsidi2}
```

# gcc-14 updates: New pass f-m-o (**F**old **M**emory **O**ffsets)

> Why do fwprop or combine not what you want to do?

- Handles Partially Overlapping memory offset calculations.
- Handle multiple memory operations sharing some intermediate calculations.
- Handle late (post Reg Alloc) offset calculation (Stack pointer)

2023-10-16 04c9cf5c786b Implement new RTL optimizations pass: fold-mem-offsets

```
-ffold-mem-offsets

add t4, sp, 16      add t2, a6, sp
add t2, a6, t4
shl t3, t2, 1
ld a2, 0(t3)        shl t3, t2, 1
add a2, 1           ld a2, 32(t3) # 16 << 1
sd a2, 8(t2)        add a2, 1
                   sd a2, 24(t2) # 8 + 16
```

# Is gcc-14 worth it: how to quantify?

**In an ideal world one could run `linux perf` on SPEC workloads and get aggregate cycle and instructions counts as well as statistical hotspots.**

- Assumes working *perf tooling* on available *hardware* to execute *fast enough*.
- All three of which are kinda questionable in RISC-V ecosystem ATM.

**Instead we measure dynamic icounts (vs. cycles) on Qemu user (vs. hardware) to run RISC-V binaries on an x86 host (fast enough).**

- Kind of approximation but good first order
- Keeps development uarch agnostic and Vendor Neutral

**Also this is a *SPEC centric world view*: one needs to start somewhere.**

- To be clear this is not “best toggle for last perf drop drag race” (yet)
- The changes are fairly widely applicable to common workloads
  - Linux kernel build

# SPEC dynamic icounts: RISC-V gcc-14 vs. gcc-13


**Build:** -Ofast -flto -static (no other toggle gimmicks)

**Run:** QEMU user with icount instrumentation plugin

**Summary:** 37,858,380,268,831 vs. 51,604,882,528,463 (~20% improvement, SMALLer is better)

SPEC2017 icounts: gcc 14.1 vs. gcc 13.1

(smaller is better)




# gcc-14 SPEC icounts: RISC-V vs. aarch64

Build: -Ofast -flto -static

- Exact same gcc-14.1 release sources based build - from scratch
- Numbers are for engineering, not marketing ;-)

## Clearly much needs to be done

SPEC icounts: RISC-V vs. aarch64 gcc 14.1 based



	aarch64 build of gcc 14.1	risc-v build of gcc 14.1
	-march=armv9-a+sve2	rv64gcv_zba_zbb_zbs_zicond_zfa
	-Ofast -flto=auto -static -ftrunc-vec-ctor	-Ofast -flto=auto -static -ftrunc-vec-ctor
	SVE2 codegen	Autovec enabled
benchmark	icount	icount vs. SVE2
500.perlbench_r-0	1,235,624,734,221	1,214,165,806,187 1.74%
500.perlbench_r-1	742,891,581,390	740,612,659,468 0.31%
500.perlbench_r-2	661,146,513,951	692,130,179,205 -4.69%
502.gcc_r-0	176,331,060,132	190,820,190,225 -8.22%
502.gcc_r-1	207,728,913,140	225,747,768,792 -8.67%
502.gcc_r-2	211,762,914,774	220,369,866,376 -4.06%
502.gcc_r-3	165,293,827,609	179,112,354,086 -8.36%
502.gcc_r-4	197,915,967,554	219,301,730,250 -10.81%
503.bwaves_r-0	198,852,846,746	278,417,236,845 -40.01%
503.bwaves_r-1	314,336,634,073	441,883,878,566 -40.58%
503.bwaves_r-2	245,770,033,150	343,713,501,705 -39.85%
503.bwaves_r-3	297,478,228,683	417,095,362,533 -40.21%
505.mcf_r	708,341,720,281	669,319,257,408 5.51%
507.cactiBSSN_r	1,363,217,761,602	2,658,677,279,673 -95.03%
508.namd_r	1,656,081,958,836	1,852,840,216,081 -11.88%
510.parest_r	1,316,493,618,361	1,660,144,903,127 -26.10%
511.povray_r	2,671,431,025,557	2,968,178,234,482 -11.11%
519.lbm_r	1,127,988,592,043	1,158,281,998,523 -2.69%
520.omnetpp_r	992,624,761,027	1,021,753,891,939 -2.93%
521.wrf_r	1,254,789,941,546	1,694,129,588,972 -35.01%
523.xalancbmk_r	1,033,436,287,173	849,247,863,343 17.82%
525.x264_r-0	201,189,068,056	277,816,418,721 -38.09%
525.x264_r-1	575,132,393,950	927,291,038,634 -61.23%
525.x264_r-2	579,805,512,651	915,363,463,243 -57.87%
526.blender_r	1,390,927,746,049	1,646,638,692,033 -18.38%
527.cam4_r	1,390,278,884,417	1,477,704,738,785 -6.29%
531.deepsjeng_r	1,494,598,201,273	1,641,969,532,631 -9.86%
538.imagick_r	2,711,178,166,664	1,912,470,009,980 29.46%
541.leela_r	1,746,137,831,038	1,983,551,291,459 -13.60%
544.nab_r	1,476,182,421,522	1,497,195,716,980 -1.42%
548.exchange2_r	2,143,325,962,206	2,072,601,275,681 3.30%
549.fotonik3d_r	556,667,724,949	1,001,472,304,174 -79.90%
554.roms_r	593,667,196,141	1,028,056,051,814 -73.17%
557.xz_r-0	342,409,623,386	363,827,039,439 -6.25%
557.xz_r-1	898,492,000,057	906,649,112,357 -0.91%
557.xz_r-2	487,117,825,355	509,023,897,942 -4.50%
997.speccrand_fr	406,197,751	402,958,586 0.80%
999.speccrand_lr	406,197,751	402,958,586 0.80%
	33,367,461,875,065	37,858,380,268,831 115.74%

# gcc-15 and beyond

## Continued support for even more Extensions

- Zbkb - Scalar Crypto but has some Zbb like insns
  - PACK is really interesting for constant synthesis

## New Pass for Enhanced Sign Extension Elimination

- Tracks register updates at byte, half-word and word-level

## SATD support in the middle-end Auto-Vectorizer

## Stack/Array Access codegen improvements


2024-03-06 9926c40a902e RISC-V: avoid LUI based const mat in alloca epilogue expansion

2024-05-13 f9cfc192ed01 RISC-V: avoid LUI based const mat in prologue/epilogue expansion [PR/105733]

2024-05-13 4bfc4585c993 RISC-V: avoid LUI based const materialization ... [part of PR/106265]

## Eliding excessive Stack Spills during instruction scheduling

<https://fprox.substack.com/p/risc-v-scalar-bit-manipulation-extensions>



# gcc @ Rivos: Community Engagement

Top 20 contributors in gcc-14 cycle

Very Active in the RISE Initiative

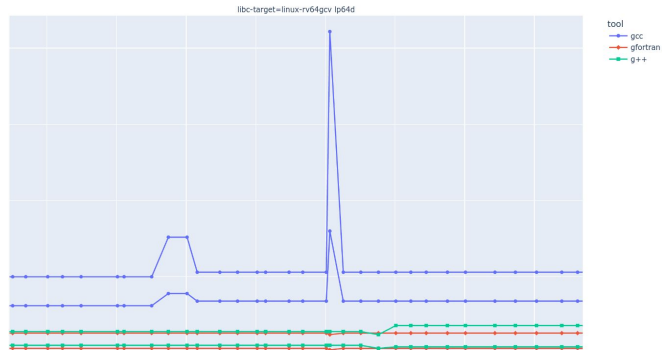
Setup public github CI for [pre-commit](#) (ewlu@rivosins.com) sanity check and [post-commit](#) (patrick@rivosinc.com).

Setup [Fuzzer](#) to find latent bugs.

Detailed Analysis of RISC-V SPEC codegen quality vs. aarch64.

Helped drive RISE 3rd party contracts for further gcc improvements.

```
1. 1876 redhat.com>
2. 801 arm.com>
3. 785 adacore.com>
4. 765 gmail.com>
5. 571 embecosm.com>
6. 545 rivai.ai>
7. 512 intel.com>
8. 449 suse.de>
9. 342 gcc.gnu.org>
10. 294 codesourcery.com>
11. 190 loongson.cn>
12. 182 linux.ibm.com>
13. 160 marvell.com>
14. 152 ventanamicro.com>
15. 152 suse.cz>
16. 140 googlemail.com>
17. 112 rivosinc.com>
18. 105 nextmovesoftware.com>
19. 95 linaro.org>
20. 93 oracle.com>
21. 87 sandoe.co.uk>
22. 85 xrylll.site>
23. 84 gjay.de>
24. 79 eswincomputing.com>
25. 71 gmx.de>
26. 71 baylibre.com>
27. 66 gdcproject.org>
28. 59 jakubdupak.com>
29. 57 CeBiTec.Uni-Bielefeld.DE>
30. 55 sifive.com>
--
```





# Thank You !

[vineetg@rivosinc.com](mailto:vineetg@rivosinc.com)

# Rivos gcc community Engagement: Competitive Analysis

## Detailed analysis of RISC-V SPEC codegen (vs. aarch64)

- Truly apples-apples : nix based workflow to use exact same gcc sources to build gcc, build SPEC and run workloads in same test harness
- Helped drive RISE Technical proposals and RFP for 3rd party contractors for further gcc improvements
- Expedites gcc improvement needed in software ecosystem